



Adding a new controller module to each cluster

ONTAP MetroCluster

NetApp
April 28, 2021

This PDF was generated from https://docs.netapp.com/us-en/ontap-metrocluster/upgrade/task_prepare_for_the_upgrade_add_2nd_controller_to_create_ha_pair.html on April 28, 2021. Always check docs.netapp.com for the latest.

Table of Contents

Adding a new controller module to each cluster	1
Preparing for the upgrade	2
Clearing the configuration on a controller module	3
Preparing cluster ports on an existing controller module	3
Preparing the netboot server to download the image	7
Setting the HA mode on the existing controller module	8
Shutting down the existing controller module	8
Installing and cabling the new controller module	9
Powering up both controller modules and displaying the LOADER prompt	11
Changing the ha-config setting on the existing and new controller modules	12
Setting the partner system ID for both controller modules	13
Booting the existing controller module	13
Assigning disks to the new controller module	13
Netbooting and setting up ONTAP on the new controller module	14
Mirroring the root aggregate on the new controller	18
Configuring intercluster LIFs	19
Creating a mirrored data aggregate on each node	26
Installing licenses for the new controller module	27
Creating unmirrored data aggregates	28
Installing the firmware after adding a controller module	30

Adding a new controller module to each cluster

You must add a new controller module to each site, creating an HA pair in each site. This is a multistep process involving both hardware and software changes that must be performed in the proper order at each site.

- The new controller module must be received from NetApp as part of the upgrade kit.

You should verify that PCIe cards in the new controller module are compatible and supported by the new controller module.

[NetApp Hardware Universe](#)

- Your system must have an empty slot available for the new controller module when upgrading to a single-chassis HA pair (an HA pair in which both controller modules reside in the same chassis).



This configuration is not supported on all systems. Platforms with single chassis configurations that are supported in ONTAP 9 are AFF A300, FAS8200, FAS8300, AFF A400, AFF80xx, FAS8020, FAS8060, FAS8080, and FAS9000.

- You must have rack space and cables for the new controller module when upgrading to a dual-chassis HA pair (an HA pair in which the controller modules reside in separate chassis).



This configuration is not supported on all systems.

- You must connect each controller module to the management network through its e0a port or, if your system has one, you can connect to the e0M port as the management port.
- These tasks must be repeated at each site.
- The preexisting controller modules are referred to as the *existing* controller modules.

The examples in this procedure have the console prompt `existing_ctlr>`.

- The controller modules that are being added are referred to as the *new* controller modules; the examples in this procedure have the console prompt `new_ctlr>`.
- This task uses the following workflow:



Preparing for the upgrade

Before upgrading to an HA pair, you must verify that your system meets all requirements and that you have all of the necessary information.

Steps

1. You need to identify unassigned disks or spare disks with available partitions that you can assign to the new controller module.

[Physical Storage Management Guide](#)

[Disk and aggregate management](#)

1. Based on the results of the previous step, perform either of the following:

If the result showed...	Then...
Not enough spare disks available for the new controller module on a system without root-data partitioning	Contact technical support for more information.
	<p>Complete the following substeps:</p> <ol style="list-style-type: none"> Determine where the aggregates for the existing node are located: <code>storage aggregate show</code> If disk ownership automatic assignment is on, turn it off: <code>storage disk option modify -node node_name -autoassign off</code> Remove ownership on disks that do not have aggregates on them: <code>storage disk removeowner disk_name</code> Repeat the previous step for as many disks as you need for the new node.

2. Verify that you have cables ready for the following connections:

- Cluster connections

If you are creating a two-node switchless cluster, you require two cables to connect the controller modules. Otherwise, you require a minimum of four cables, two for each controller module connection to the cluster-network switch. Other systems (like the 80xx series) have defaults of either four or six cluster connections.

- HA interconnect connections, if the system is in a dual-chassis HA pair

3. Verify that you have a serial port console available for the controller modules.

4. Verify that your environment meets the site and system requirements.

[NetApp Hardware Universe](#)

5. Gather all of the IP addresses and other network parameters for the new controller module.

Clearing the configuration on a controller module

Unresolved directive in <stdin> - include:::../_include/steps_clear_config_boot_to_maint_mode.adoc[]

Preparing cluster ports on an existing controller module

Before installing a new controller module, you must configure cluster ports on the existing controller module so that the cluster ports can provide cluster communication with the new controller module.

If you are creating a two-node switchless cluster (with no cluster network switches), you must enable the

switchless cluster networking mode.

For detailed information about port, LIF, and network configuration in ONTAP, see the [Network Management Guide](#).

Steps

1. Determine which ports should be used as the node's cluster ports.

For a list of the default port roles for your platform, see the [Hardware Universe](#)

The *Installation and Setup Instructions* for your platform on the NetApp Support Site contains information about the ports for cluster network connections.

2. For each cluster port, identify the port roles: **network port show**

In the following example, ports e0a, e0b, e0c, and e0d must be changed to cluster ports:

```
cluster_A::> network port show
```

```
Node: controller_A_1
```

```
Speed(Mbps) Health
```

Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper	Status
e0M	Default	mgmt_bd_1500	up	1500	auto/1000	healthy
e0a	Default	Default	up	1500	auto/10000	healthy
e0b	Default	Default	up	1500	auto/10000	healthy
e0c	Default	Default	up	1500	auto/10000	healthy
e0d	Default	Default	up	1500	auto/10000	healthy
e0i	Default	Default	down	1500	auto/10	-
e0j	Default	Default	down	1500	auto/10	-
e0k	Default	Default	down	1500	auto/10	-
e0l	Default	Default	down	1500	auto/10	-
e2a	Default	Default	up	1500	auto/10000	healthy
e2b	Default	Default	up	1500	auto/10000	healthy
e4a	Default	Default	up	1500	auto/10000	healthy
e4b	Default	Default	up	1500	auto/10000	healthy

```
13 entries were displayed.
```

3. For any data LIF that is using a cluster port as the home-port or current-port, modify the LIF to use a data port as its home-port: **network interface modify**

The following example changes the home port of a data LIF to a data port:

```
cluster1::> network interface modify -lif datalif1 -vserver vs1 -home  
-port e1b
```

4. For each LIF that you modified, revert the LIF to its new home port: **network interface revert**

The following example reverts the LIF datalif1 to its new home port e1b:

```
cluster1::> network interface revert -lif datalif1 -vserver vs1
```

5. Remove any VLAN ports using cluster ports as member ports and ifgrps using cluster ports as member ports.

- a. Delete VLAN ports:

network port vlan delete -node node-name -vlan-name portid-vlandid

For example:

```
network port vlan delete -node node1 -vlan-name e1c-80
```

- b. Remove physical ports from the interface groups:

network port ifgrp remove-port -node node-name -ifgrp interface-group-name -port portid

For example:

```
network port ifgrp remove-port -node node1 -ifgrp ala -port e0d
```

- c. Remove VLAN and interface group ports from broadcast domain::

***network port broadcast-domain remove-ports -ipspace ipspace -broadcast -domain broadcast-domain-name -ports nodename:portname,nodename:portname,..**

- d. Modify interface group ports to use other physical ports as member as needed.:

***ifgrp add-port -node node-name -ifgrp interface-group-name -port port-id**

6. Verify that the port roles have changed: **network port show**

The following example shows that ports e0a, e0b, e0c, and e0d are now cluster ports:

Node: controller_A_1

Speed(Mbps) Health

Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper	Status
e0M	Default	mgmt_bd_1500	up	1500	auto/1000	healthy
e0a	Cluster	Cluster	up	9000	auto/10000	healthy
e0b	Cluster	Cluster	up	9000	auto/10000	healthy
e0c	Cluster	Cluster	up	9000	auto/10000	healthy
e0d	Cluster	Cluster	up	9000	auto/10000	healthy
e0i	Default	Default	down	1500	auto/10 -	
e0j	Default	Default	down	1500	auto/10 -	
e0k	Default	Default	down	1500	auto/10 -	
e0l	Default	Default	down	1500	auto/10 -	
e2a	Default	Default	up	1500	auto/10000	healthy
e2b	Default	Default	up	1500	auto/10000	healthy
e4a	Default	Default	up	1500	auto/10000	healthy
e4b	Default	Default	up	1500	auto/10000	healthy

13 entries were displayed.

7. If your system is part of a switched cluster, create cluster LIFs on the cluster ports: **network interface create**

The following example creates a cluster LIF on one of the node's cluster ports. The **-auto** parameter configures the LIF to use a link-local IP address.

```
cluster1::> network interface create -vserver Cluster -lif clus1 -role
cluster -home-node node0 -home-port e1a -auto true
```

8. If you are creating a two-node switchless cluster, enable the switchless cluster networking mode:

- a. Change to the advanced privilege level from either node:

set -privilege advanced

You can respond **y** when prompted whether you want to continue into advanced mode. The advanced mode prompt appears (***>**).

- b. Enable the switchless cluster networking mode: **network options switchless-cluster modify -enabled true**

- c. Return to the admin privilege level: **set -privilege admin**



Cluster interface creation for the existing node in a two-node switchless cluster system is completed after cluster setup is completed through a netboot on the new controller module.

Preparing the netboot server to download the image

When you are ready to prepare the netboot server, you must download the correct ONTAP netboot image from the NetApp Support Site to the netboot server and note the IP address.

- You must be able to access an HTTP server from the system before and after adding the new controller module.
- You must have access to the NetApp Support Site to download the necessary system files for your platform and your version of ONTAP.

[NetApp Support Site](#)


- Both controller modules in the HA pair must run the same version of ONTAP.


Steps

1. Download the appropriate ONTAP software from the software download section of the NetApp Support Site and store the `<ontap_version>_image.tgz` file on a web-accessible directory.

The `<ontap_version>_image.tgz` file is used for performing a netboot of your system.

2. Change to the web-accessible directory and verify that the files you need are available.

For...	Then...
FAS2200, FAS2500, FAS3200, FAS6200, FAS/AFF8000 series systems	<div>Extract the contents of the <code><ontap_version>_image.tgz</code> file to the target directory:</div> <div><pre>tar -zxvf <ontap_version>_image.tgz</pre></div> <div><div></div><div>If you are extracting the contents on Windows, use 7-Zip or WinRAR to extract the netboot image.</div></div> <div>Your directory listing should contain a netboot folder with a kernel file:</div> <div><pre>netboot/kernel</pre></div>

For...	Then...
All other systems	<p>Your directory listing should contain the following file:</p> <pre><ontap_version>_image.tgz</pre> <div>  <p>There is no need to extract the file contents.</p> </div>

- Determine the IP address of the existing controller module.

This address is referred to later in this procedure as *ip-address-of-existing controller*.

- Ping *ip-address-of-existing controller* to verify that the IP address is reachable.

Setting the HA mode on the existing controller module

You must use the storage failover modify command to set the mode on the existing controller module. The mode value is enabled later, after you reboot the controller module.

Steps

- Set the mode to HA:

```
storage failover modify -mode ha -node existing_node_name
```

Shutting down the existing controller module

You must perform a clean shutdown of the existing controller module to verify that all of the data has been written to disk. You must also disconnect the power supplies.

- Halt the node from the existing controller module prompt: **halt local -inhibit-takeover true**

If you are prompted to continue the halt procedure, enter **y** when prompted, and then wait until the system stops at the LOADER prompt.



You must perform a clean system shutdown before replacing the system components to avoid losing unwritten data in the NVRAM or NVMEM.

- In an 80xx system, the NVRAM LED is located on the controller module to the right of the network ports, marked with a battery symbol. This LED blinks if there is unwritten data in the NVRAM. If this LED is flashing amber after you enter the halt command, you need to reboot your system and try halting it again.
- If you are not already grounded, properly ground yourself.
 - Turn off the power supplies and disconnect the power, using the correct method for your system and

power-supply type:

If your system uses...	Then...
AC power supplies	Unplug the power cords from the power source, and then remove the power cords.
DC power supplies	Remove the power at the DC source, and then remove the DC wires, if necessary.

Installing and cabling the new controller module

You must physically install the new controller module in the chassis, and then cable it.

Steps

1. If you have an I/O expansion module (IOXM) in your system and are creating a single-chassis HA pair, you must uncable and remove the IOXM.

You can then use the empty bay for the new controller module. However, the new configuration will not have the extra I/O provided by the IOXM.

2. Physically install the new controller module and, if necessary, install additional fans:

If you are adding a controller module...	Then perform these steps...
To an empty bay to create a single-chassis HA pair and the system belongs to one of the following platforms:	<ol style="list-style-type: none">a. Remove the blank plate in the rear of the chassis that covers the empty bay that will contain the new controller module.b. Gently push the controller module halfway into the chassis. <p>To prevent the controller module from automatically booting, do not fully seat it in the chassis until later in this procedure.</p>
In a separate chassis from its HA partner to create a dual-chassis HA pair when the existing configuration is in a controller-IOX module configuration. <ul style="list-style-type: none">• FAS8200• 80xx	Install the new system in the rack or system cabinet.

3. Cable the cluster network connections, as necessary:
 - a. Identify the ports on the controller module for the cluster connections.

[AFF A320 systems: Installation and setup](#)

[AFF A220/FAS2700 Systems Installation and Setup Instructions](#)

- b. If you are configuring a switched cluster, identify the ports that you will use on the cluster network switches.

See the *Clustered Data ONTAP Switch Setup Guide for Cisco Switches*, *NetApp 10G Cluster-Mode Switch Installation Guide* or *NetApp 1G Cluster-Mode Switch Installation Guide*, depending on what switches you are using.

- c. Connect cables to the cluster ports:

If the cluster is...	Then...
A two-node switchless cluster	Directly connect the cluster ports on the existing controller module to the corresponding cluster ports on the new controller module.
A switched cluster	Connect the cluster ports on each controller to the ports on the cluster network switches identified in substep b.

Cabling the new controller module's FC-VI and HBA ports to the FC switches

The new controller module's FC-VI ports and HBAs (host bus adapters) must be cabled to the site FC switches.

Steps

1. Cable the FC-VI ports and HBA ports, using the table for your configuration and switch model.
 - [Port assignments for FC switches when using ONTAP 9.1 and later](#)
 - [Port assignments for FC switches when using ONTAP 9.0](#)
 - [Port assignments for systems using two initiator ports](#)

Cabling the new controller module's cluster peering connections

You must cable the new controller module to the cluster peering network so that it has connectivity with the cluster on the partner site.

At least two ports on each controller module should be used for cluster peering.

The recommended minimum bandwidth for the ports and network connectivity is 1 GbE.

Steps

1. Identify and cable at least two ports for cluster peering and verify they have network connectivity with the partner cluster.

Powering up both controller modules and displaying the LOADER prompt

You power up the existing controller module and the new controller module to display the LOADER prompt.

Steps

1. Power up the controller modules and interrupt the boot process, using the steps for your configuration:

If the controller modules are...	Then...
In the same chassis	<ol style="list-style-type: none">a. Verify that the new controller module is not fully inserted into the bay. The existing controller module should be fully inserted into the bay because it was never removed from the chassis, but the new controller module should not be.b. Connect the power and turn on the power supplies so that the existing controller module receives power.c. Interrupt the boot process on the existing controller module by pressing Ctrl-C.d. Push the new controller module firmly into the bay. When fully seated, the new controller module receives power and automatically boots.e. Interrupt the boot process by pressing Ctrl-C.f. Tighten the thumbscrew on the cam handle, if present.g. Install the cable management device, if present.h. Bind the cables to the cable management device with the hook and loop strap.
In separate chassis	<ol style="list-style-type: none">a. Turn on the power supplies on the existing controller module.b. Interrupt the boot process by pressing Ctrl-C.c. Repeat these steps for the new controller module

Each controller module should display the LOADER prompt (**LOADER>**, **LOADER-A>**, or **LOADER-B>**).



If there is no LOADER prompt, record the error message and contact technical support. If the system displays the boot menu, reboot and attempt to interrupt the boot process again.

Changing the ha-config setting on the existing and new controller modules

When you expand a MetroCluster configuration, you must update the ha-config setting of the existing controller module and the new controller module. You must also determine the system ID of the new controller module.

This task is performed in Maintenance mode on both the existing and new controller modules.

Steps

1. Change the ha-config setting of the existing controller module:

- a. Display the ha-config setting of the existing controller module and chassis:

```
ha-config show
```

The ha-config setting is `mcc-2n` for all components because the controller module was in a two-node MetroCluster configuration.

- b. Change the ha-config setting of the existing controller module to `mcc`:

```
ha-config modify controller mcc
```

- c. Change the ha-config setting of the existing chassis to `mcc`:

```
ha-config modify chassis mcc
```

- d. Retrieve the system ID for the existing controller module:

```
sysconfig
```

Note the system ID. You need it when you set the partner ID on the new controller module.

- e. Exit Maintenance mode to return to the LOADER prompt:

```
halt
```

2. Change the ha-config setting and retrieve the system ID of the new controller module:

- a. If the new controller module is not already in Maintenance mode, boot it to Maintenance mode:

```
boot_ontap maint
```

- b. Change the ha-config setting of the new controller module to `mcc`:

```
ha-config modify controller mcc
```

- c. Change the ha-config setting of the new chassis to `mcc`:

```
ha-config modify chassis mcc
```

- d. Retrieve the system ID for the new controller module:

```
sysconfig
```

Note the system ID. You need it when you set the partner ID and assign disks to the new controller module.

- e. Exit Maintenance mode to return to the LOADER prompt:

```
halt
```

Setting the partner system ID for both controller modules

You must set the partner system ID on both controller modules so that they can form an HA pair.

This task is performed with both controller modules at the LOADER prompt.

Steps

1. On the existing controller module, set the partner system ID to that of the new controller module:

```
setenv partner-sysid sysID_of_new_controller
```

2. On the new controller module, set the partner system ID to that of the existing controller module:

```
setenv partner-sysid sysID_of_existing_controller
```

Booting the existing controller module

You must boot the existing controller module to ONTAP.

Steps

1. At the LOADER prompt, boot the existing controller module to ONTAP:

```
boot_ontap
```

Assigning disks to the new controller module

Before you complete the configuration of the new controller module through netboot, you must assign disks to it.

You must have made sure that there are enough spares, unassigned disks, or assigned disks that are not part of an existing aggregate.

Preparing for the upgrade

These steps are performed on the existing controller module.

Steps

1. Assign the root disk to the new controller module:

```
storage disk assign -disk disk_name -sysid new_controller_sysID -force true
```

If your platform model uses the Advanced Drive Partitioning (ADP) feature, you must include the `-root true` parameter:

```
storage disk assign -disk disk_name -root true -sysid new_controller_sysID
```

`-force true`

2. Assign the remaining required disks to the new controller module by entering the following command for each disk:

```
storage disk assign -disk disk_name -sysid new_controller_sysID -force true
```

3. Verify that the disk assignments are correct:

```
storage disk show -partitionownership
```



Ensure that you have assigned all disks that you intend to assign to the new node.

Netbooting and setting up ONTAP on the new controller module

You must perform a specific sequence of steps to netboot and install the ONTAP operating system on the new controller module when adding controller modules to an existing MetroCluster configuration.

- This task starts at the LOADER prompt of the new controller module.
- This task includes initializing disks.

The amount of time you need to initialize the disks depends on the size of the disks.


- The system automatically assigns two disks to the new controller module.

[Disk and aggregate management](#)

Steps

1. At the LOADER prompt, configure the IP address of the new controller module based on DHCP availability:

If DHCP is...	Then enter the following command...
Available	<code>ifconfig e0M -auto</code>

If DHCP is...	Then enter the following command...
Not available	<pre>ifconfig e0M -addr=filer_addr - mask=netmask -gw=gateway -dns=dns_addr -domain=dns_domain</pre> <p><code>filer_addr</code> is the IP address of the storage system.</p> <p><code>netmask</code> is the network mask of the storage system.</p> <p><code>gateway</code> is the gateway for the storage system.</p> <p><code>dns_addr</code> is the IP address of a name server on your network.</p> <p><code>dns_domain</code> is the Domain Name System (DNS) domain name. If you use this optional parameter, you do not need a fully qualified domain name in the netboot server URL; you need only the server's host name.</p> <div>  <p>Other parameters might be necessary for your interface. For details, use the <code>help ifconfig</code> command at the LOADER prompt.</p> </div>

2. At the LOADER prompt, netboot the new node:

For...	Then...
FAS2200, FAS2500, FAS3200, FAS6200, FAS/AFF8000 series systems	<pre>netboot http://web_server_ip/path_to_web- accessible_directory/netboot/kernel</pre>
All other systems	<pre>netboot http://web_server_ip/path_to_web- accessible_directory/<ontap_version>_i mage.tgz</pre>

The `path_to_the_web-accessible_directory` is the location of the downloaded `<ontap_version>_image.tgz` file.

1. Select the **Install new software first** option from the displayed menu.

This menu option downloads and installs the new ONTAP image to the boot device.

- You should enter **y** when prompted with the message that this procedure is not supported for

nondisruptive upgrade on an HA pair.

- You should enter **y** when warned that this process replaces the existing ONTAP software with new software.
- You should enter the path as follows when prompted for the URL of the image.tgz file: + **http://path_to_the_web-accessible_directory/image.tgz**

2. Enter **y** when prompted regarding nondisruptive upgrade or replacement of the software.

3. Enter the path to the image.tgz file when prompted for the URL of the package.

```
What is the URL for the package? `http://path_to_web-
accessible_directory/image.tgz`
```

4. Enter **n** to skip the backup recovery when prompted to restore the backup configuration.

```
*****
*                               *
*           Restore Backup Configuration           *
* This procedure only applies to storage controllers that *
* are configured as an HA pair.                      *
*                               *
* Choose Yes to restore the "varfs" backup configuration *
* from the SSH server. Refer to the Boot Device Replacement *
* guide for more details.                             *
* Choose No to skip the backup recovery and return to the *
* boot menu.                                           *
*****

Do you want to restore the backup configuration
now? {y|n} `n`
```

5. Enter **y** when prompted to reboot now.

```
The node must be rebooted to start using the newly installed software.
Do you want to
reboot now? {y|n} `y`
```

6. If necessary, select the option to **Clean configuration and initialize all disks** after the node has booted.

Because you are configuring a new controller module and the new controller module's disks are empty, you can respond **y** when the system warns you that this will erase all disks.



The amount of time needed to initialize disks depends on the size of your disks and configuration.

7. After the disks are initialized and the Cluster Setup wizard starts, set up the node:

- a. Enter the node management LIF information on the console.
8. Log in to the node, and enter the **cluster setup** and then enter **join** when prompted to join the cluster.

```
Do you want to create a new cluster or join an existing cluster?
{create, join}: `join`
```

9. Respond to the remaining prompts as appropriate for your site.

The [Software Setup Guide](#) for your version of ONTAP contains additional details.

10. If the system is in a two-node switchless cluster configuration, create the cluster interfaces on the existing node using the network interface create command to create cluster LIFs on the cluster ports.

The following is an example command for creating a cluster LIF on one of the node's cluster ports. The **-auto** parameter configures the LIF to use a link-local IP address.

```
cluster_A::> network interface create -vserver Cluster -lif clus1 -role
cluster -home-node node_A_1 -home-port e1a -auto true
```

11. After setup is complete, verify that the node is healthy and eligible to participate in the cluster:

cluster show

The following example shows a cluster after the second node (cluster1-02) has been joined to it:

```
cluster_A::> cluster show
Node                               Health  Eligibility
-----
node_A_1                           true    true
node_A_2                           true    true
```

You can access the Cluster Setup wizard to change any of the values you entered for the admin storage virtual machine (SVM) or node SVM by using the cluster setup command.

12. Confirm that you have four ports configured as cluster interconnects:

network port show

The following example shows output for two controller modules in cluster_A:

```
cluster_A::> network port show
```

(Mbps)					Speed	
Node	Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper

node_A_1						
	**e0a	Cluster	Cluster	up	9000	
	auto/1000					
	e0b	Cluster	Cluster	up	9000	
	auto/1000**					
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
	e0e	Default	Default	up	1500	auto/1000
	e0f	Default	Default	up	1500	auto/1000
	e0g	Default	Default	up	1500	auto/1000
node_A_2						
	**e0a	Cluster	Cluster	up	9000	
	auto/1000					
	e0b	Cluster	Cluster	up	9000	
	auto/1000**					
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
	e0e	Default	Default	up	1500	auto/1000
	e0f	Default	Default	up	1500	auto/1000
	e0g	Default	Default	up	1500	auto/1000

14 entries were displayed.

Mirroring the root aggregate on the new controller

You must mirror the root aggregate to provide data protection when you are adding a controller to a MetroCluster configuration.

This task must be performed on the new controller module.

1. Mirror the root aggregate:

```
storage aggregate mirror aggr_name
```

The following command mirrors the root aggregate for controller_A_1:

```
controller_A_1::> storage aggregate mirror aggr0_controller_A_1
```

This mirrors the aggregate, so it consists of a local plex and a remote plex located at the remote MetroCluster site.

Configuring intercluster LIFs

You must create intercluster LIFs on ports used for communication between the MetroCluster partner clusters. You can use dedicated ports or ports that also have data traffic.

Configuring intercluster LIFs on dedicated ports

You can configure intercluster LIFs on dedicated ports. Doing so typically increases the available bandwidth for replication traffic.

Steps

1. List the ports in the cluster:

```
network port show
```

For complete command syntax, see the man page.

The following example shows the network ports in cluster01:

```
cluster01::> network port show
```

						Speed
(Mbps)						
Node	Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper
-----	-----	-----	-----	-----	-----	-----
cluster01-01						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
	e0e	Default	Default	up	1500	auto/1000
	e0f	Default	Default	up	1500	auto/1000
cluster01-02						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
	e0e	Default	Default	up	1500	auto/1000
	e0f	Default	Default	up	1500	auto/1000

2. Determine which ports are available to dedicate to intercluster communication:

```
network interface show -fields home-port,curr-port
```

For complete command syntax, see the man page.

The following example shows that ports e0e and e0f have not been assigned LIFs:

```
cluster01::> network interface show -fields home-port,curr-port
vserver lif                home-port curr-port
-----
Cluster cluster01-01_clus1  e0a      e0a
Cluster cluster01-01_clus2  e0b      e0b
Cluster cluster01-02_clus1  e0a      e0a
Cluster cluster01-02_clus2  e0b      e0b
cluster01
      cluster_mgmt          e0c      e0c
cluster01
      cluster01-01_mgmt1    e0c      e0c
cluster01
      cluster01-02_mgmt1    e0c      e0c
```

3. Create a failover group for the dedicated ports:

```
network interface failover-groups create -vserver system_SVM -failover-group  
failover_group -targets physical_or_logical_ports
```

The following example assigns ports e0e and e0f to the failover group intercluster01 on the system SVMcluster01:

```
cluster01::> network interface failover-groups create -vserver cluster01  
-failover-group  
intercluster01 -targets  
cluster01-01:e0e,cluster01-01:e0f,cluster01-02:e0e,cluster01-02:e0f
```

4. Verify that the failover group was created:

```
network interface failover-groups show
```

For complete command syntax, see the man page.

```

cluster01::> network interface failover-groups show

```

Vserver	Group	Failover Targets
Cluster	Cluster	cluster01-01:e0a, cluster01-01:e0b, cluster01-02:e0a, cluster01-02:e0b
cluster01	Default	cluster01-01:e0c, cluster01-01:e0d, cluster01-02:e0c, cluster01-02:e0d, cluster01-01:e0e, cluster01-01:e0f cluster01-02:e0e, cluster01-02:e0f
	intercluster01	cluster01-01:e0e, cluster01-01:e0f cluster01-02:e0e, cluster01-02:e0f

5. Create intercluster LIFs on the system SVM and assign them to the failover group.

ONTAP version	Command
9.6 and later	<code>network interface create -vserver system_SVM -lif LIF_name -service-policy default-intercluster -home-node node -home -port port -address port_IP -netmask netmask -failover -group failover_group</code>
9.5 and earlier	<code>network interface create -vserver system_SVM -lif LIF_name -role intercluster -home-node node -home-port port -address port_IP -netmask netmask -failover-group failover_group</code>

For complete command syntax, see the man page.

The following example creates intercluster LIFs `cluster01_icl01` and `cluster01_icl02` in the failover group `intercluster01`:

```

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0e
-address 192.168.1.201
-netmask 255.255.255.0 -failover-group intercluster01

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0e
-address 192.168.1.202
-netmask 255.255.255.0 -failover-group intercluster01

```

6. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:

```
network interface show -service-policy default-intercluster
```

In ONTAP 9.5 and earlier:

```
network interface show -role intercluster
```

For complete command syntax, see the man page.

```

cluster01::> network interface show -service-policy default-intercluster

```

Current Is	Logical	Status	Network	Current
Vserver	Interface	Admin/Oper	Address/Mask	Node
Home				Port
cluster01	cluster01_icl01	up/up	192.168.1.201/24	cluster01-01 e0e
true	cluster01_icl02	up/up	192.168.1.202/24	cluster01-02 e0f
true				

7. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:


```
network interface show -service-policy default-intercluster -failover
```

In ONTAP 9.5 and earlier:

```
network interface show -role intercluster -failover
```

For complete command syntax, see the man page.

The following example shows that the intercluster LIFs `cluster01_icl01` and `cluster01_icl02` on the SVMe0e port will fail over to the e0f port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
```

Vserver	Logical Interface	Home Node:Port	Failover Policy	Failover Group
cluster01	cluster01_icl01	cluster01-01:e0e	local-only	
intercluster01			Failover Targets: cluster01-01:e0e, cluster01-01:e0f	
cluster01	cluster01_icl02	cluster01-02:e0e	local-only	
intercluster01			Failover Targets: cluster01-02:e0e, cluster01-02:e0f	

Configuring intercluster LIFs on shared data ports

You can configure intercluster LIFs on ports shared with the data network. Doing so reduces the number of ports you need for intercluster networking.

Steps

1. List the ports in the cluster:

```
network port show
```

For complete command syntax, see the man page.

The following example shows the network ports in cluster01:

```
cluster01::> network port show
```

(Mbps)						Speed
Node	Port	IPspace	Broadcast Domain	Link	MTU	Admin/Oper
-----	-----	-----	-----	-----	-----	

cluster01-01						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000
cluster01-02						
	e0a	Cluster	Cluster	up	1500	auto/1000
	e0b	Cluster	Cluster	up	1500	auto/1000
	e0c	Default	Default	up	1500	auto/1000
	e0d	Default	Default	up	1500	auto/1000

2. Create intercluster LIFs on the system SVM:

In ONTAP 9.6 and later:

```
network interface create -vserver system_SVM -lif LIF_name -service-policy
default-intercluster -home-node node -home-port port -address port_IP -netmask
netmask
```

In ONTAP 9.5 and earlier:

```
network interface create -vserver system_SVM -lif LIF_name -role intercluster
-home-node node -home-port port -address port_IP -netmask *netmask
```

For complete command syntax, see the man page.

The following example creates intercluster LIFs `cluster01_icl01` and `cluster01_icl02`:

```
cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl01 -service-
policy default-intercluster -home-node cluster01-01 -home-port e0c
-address 192.168.1.201
-netmask 255.255.255.0

cluster01::> network interface create -vserver cluster01 -lif
cluster01_icl02 -service-
policy default-intercluster -home-node cluster01-02 -home-port e0c
-address 192.168.1.202
-netmask 255.255.255.0
```

3. Verify that the intercluster LIFs were created:

In ONTAP 9.6 and later:

```
network interface show -service-policy default-intercluster
```

In ONTAP 9.5 and earlier:

```
network interface show -role intercluster
```

For complete command syntax, see the man page.

```
cluster01::> network interface show -service-policy default-intercluster
```

	Logical	Status	Network	Current	
Current Is					
Vserver	Interface	Admin/Oper	Address/Mask	Node	Port
Home					
-----	-----	-----	-----	-----	
-----	----				
cluster01	cluster01_icl01				
		up/up	192.168.1.201/24	cluster01-01	e0c
true					
	cluster01_icl02				
		up/up	192.168.1.202/24	cluster01-02	e0c
true					

4. Verify that the intercluster LIFs are redundant:

In ONTAP 9.6 and later:

```
network interface show -service-policy default-intercluster -failover
```

In ONTAP 9.6 and later:

In ONTAP 9.5 and earlier:

```
network interface show -role intercluster -failover
```

For complete command syntax, see the man page.

The following example shows that the intercluster LIFs `cluster01_icl01` and `cluster01_icl02` on the `e0c` port will fail over to the `e0d` port.

```
cluster01::> network interface show -service-policy default-intercluster
-failover
```

Vserver	Logical Interface	Home Node:Port	Failover Policy	Failover Group

cluster01				
	cluster01_icl01	cluster01-01:e0c	local-only	
	192.168.1.201/24			
		Failover Targets: cluster01-01:e0c,		
		cluster01-01:e0d		
	cluster01_icl02	cluster01-02:e0c	local-only	
	192.168.1.201/24			
		Failover Targets: cluster01-02:e0c,		
		cluster01-02:e0d		

Creating a mirrored data aggregate on each node

You must create a mirrored data aggregate on each node in the DR group.

- You should know what drives will be used in the new aggregate.
- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can ensure that the correct drive type is selected.
- Drives are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.

In systems using ADP, aggregates are created using partitions in which each drive is partitioned in to P1, P2 and P3 partitions.

- Aggregate names should conform to the naming scheme you determined when you planned your MetroCluster configuration.

Disk and aggregate management

1. Display a list of available spares:

```
storage disk show -spare -owner node_name
```

2. Create the aggregate by using the `storage aggregate create -mirror true` command.

If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To ensure that the aggregate is created on a specific node, use the `-node` parameter or specify drives that are owned by that node.

You can specify the following options:

- Aggregate's home node (that is, the node that owns the aggregate in normal operation)
- List of specific drives that are to be added to the aggregate
- Number of drives to include



In the minimum supported configuration, in which a limited number of drives are available, you must use the `force-small-aggregate` option to allow the creation of a three disk RAID-DP aggregate.

- Checksum style to use for the aggregate
- Type of drives to use
- Size of drives to use
- Drive speed to use
- RAID type for RAID groups on the aggregate
- Maximum number of drives that can be included in a RAID group
- Whether drives with different RPM are allowed For more information about these options, see the `storage aggregate create` man page.

The following command creates a mirrored aggregate with 10 disks:

+

```
cluster_A::> storage aggregate create aggr1_node_A_1 -diskcount 10 -node
node_A_1 -mirror true
[Job 15] Job is queued: Create aggr1_node_A_1.
[Job 15] The job is starting.
[Job 15] Job succeeded: DONE
```

1. Verify the RAID group and drives of your new aggregate:

`storage aggregate show-status -aggregate aggregate-name`

Installing licenses for the new controller module

You must add licenses for the new controller module for any ONTAP services that require standard (node-locked) licenses. For features with standard licenses, each node in the cluster must have its own key for the feature.

For detailed information about licensing, see the knowledgebase article 3013749: Data ONTAP 8.2 Licensing

Overview and References on the NetApp Support Site and the *System Administration Reference*.

Steps

1. If necessary, obtain license keys for the new node on the NetApp Support Site in the My Support section under Software licenses.

If the site does not have the license keys you need, contact your sales or support representative.

2. Issue the following command to install each license key:

```
system license add -license-code license_key
```

The *license_key* is 28 digits in length.

Repeat this step for each required standard (node-locked) license.

Creating unmirrored data aggregates

You can optionally create unmirrored data aggregates for data that does not require the redundant mirroring provided by MetroCluster configurations.

- You should know what drives or array LUNs will be used in the new aggregate.
- If you have multiple drive types in your system (heterogeneous storage), you should understand how you can verify that the correct drive type is selected.

IMPORTANT:

In MetroCluster IP configurations, remote unmirrored aggregates are not accessible after a switchover



The unmirrored aggregates must be local to the node owning them.

- Drives and array LUNs are owned by a specific node; when you create an aggregate, all drives in that aggregate must be owned by the same node, which becomes the home node for that aggregate.
- Aggregate names should conform to the naming scheme you determined when you planned your MetroCluster configuration.
- The *Disks and Aggregates Power Guide* contains more information about mirroring aggregates.

1. Enable unmirrored aggregate deployment:

```
metrocluster modify -enable-unmirrored-aggr-deployment true
```

2. Verify that disk autoassignment is disabled:

```
disk option show
```

3. Install and cable the disk shelves that will contain the unmirrored aggregates.

You can use the procedures in the Installation and Setup documentation for your platform and disk shelves.

[AFF and FAS Documentation Center](#)

4. Manually assign all disks on the new shelf to the appropriate node:

```
disk assign -disk disk-id -owner owner-node-name
```

5. Create the aggregate:

```
storage aggregate create
```

If you are logged in to the cluster on the cluster management interface, you can create an aggregate on any node in the cluster. To verify that the aggregate is created on a specific node, you should use the `-node` parameter or specify drives that are owned by that node.

You must also ensure that you are only including drives on the unmirrored shelf to the aggregate.

You can specify the following options:

- Aggregate's home node (that is, the node that owns the aggregate in normal operation)
- List of specific drives or array LUNs that are to be added to the aggregate
- Number of drives to include
- Checksum style to use for the aggregate
- Type of drives to use
- Size of drives to use
- Drive speed to use
- RAID type for RAID groups on the aggregate
- Maximum number of drives or array LUNs that can be included in a RAID group
- Whether drives with different RPM are allowed For more information about these options, see the `storage aggregate create` man page.

The following command creates a unmirrored aggregate with 10 disks:

+

```
controller_A_1:> storage aggregate create aggr1_controller_A_1
-diskcount 10 -node controller_A_1
[Job 15] Job is queued: Create aggr1_controller_A_1.
[Job 15] The job is starting.
[Job 15] Job succeeded: DONE
```

1. Verify the RAID group and drives of your new aggregate:

```
storage aggregate show-status -aggregate aggregate-name
```

2. Disable unmirrored aggregate deployment:

```
metrocluster modify -enable-unmirrored-aggr-deployment false
```

3. Verify that disk autoassignment is enabled:

`disk option show`

Related information

[Disk and aggregate management](#)

Installing the firmware after adding a controller module

After adding the controller module, you must install the latest firmware on the new controller module so that the controller module functions properly with ONTAP.

Steps

1. Download the most current version of firmware for your system and follow the instructions for downloading and installing the new firmware.

[NetApp Downloads: System Firmware and Diagnostics](#)

Copyright Information

Copyright © 2021 NetApp, Inc. All rights reserved. Printed in the U.S. No part of this document covered by copyright may be reproduced in any form or by any means-graphic, electronic, or mechanical, including photocopying, recording, taping, or storage in an electronic retrieval system- without prior written permission of the copyright owner.

Software derived from copyrighted NetApp material is subject to the following license and disclaimer:

THIS SOFTWARE IS PROVIDED BY NETAPP “AS IS” AND WITHOUT ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE, WHICH ARE HEREBY DISCLAIMED. IN NO EVENT SHALL NETAPP BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

NetApp reserves the right to change any products described herein at any time, and without notice. NetApp assumes no responsibility or liability arising from the use of products described herein, except as expressly agreed to in writing by NetApp. The use or purchase of this product does not convey a license under any patent rights, trademark rights, or any other intellectual property rights of NetApp.

The product described in this manual may be protected by one or more U.S. patents, foreign patents, or pending applications.

RESTRICTED RIGHTS LEGEND: Use, duplication, or disclosure by the government is subject to restrictions as set forth in subparagraph (c)(1)(ii) of the Rights in Technical Data and Computer Software clause at DFARS 252.277-7103 (October 1988) and FAR 52-227-19 (June 1987).

Trademark Information

NETAPP, the NETAPP logo, and the marks listed at <http://www.netapp.com/TM> are trademarks of NetApp, Inc. Other company and product names may be trademarks of their respective owners.